

Homework 3

Rohan Thakur, Charles Kekeh and Megan Jasek

February 13, 2016

```
# Load the dataframe
load("twoyear.RData")
desc
```

```
##      variable                                label
## 1   female                                =1 if female
## 2   phsrank  % high school rank; 100 = best
## 3     BA                                =1 if Bachelor's degree
## 4     AA                                =1 if Associate's degree
## 5   black                                =1 if African-American
## 6 hispanic                                =1 if Hispanic
## 7     id                                ID Number
## 8   exper  total (actual) work experience
## 9     jc                                total 2-year credits
## 10    univ                                total 4-year credits
## 11   lwage                                log hourly wage
## 12 stotal  total standardized test score
## 13 smcity                                =1 if small city, 1972
## 14 medcity                                =1 if med. city, 1972
## 15 submed  =1 if suburb med. city, 1972
## 16 lgcity                                =1 if large city, 1972
## 17 sublg   =1 if suburb large city, 1972
## 18 vlgcity =1 if very large city, 1972
## 19 subvlg  =1 if sub. very lge. city, 1972
## 20     ne                                =1 if northeast
## 21     nc                                =1 if north central
## 22    south                                =1 if south
## 23 totcoll                                jc + univ
```

Question 1


```
summary(data$lwage)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.5555  1.9250  2.2760  2.2480  2.5970  3.9120
```

```
print(quantile(data$lwage, probs = c(0.01, 0.05, 0.1,
  0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%      5%      10%      25%      50%      75%      90%      95%
## 1.148702 1.398129 1.609438 1.925291 2.276300 2.596916 2.851921 2.995732
##      99%     100%
## 3.325316 3.911953
```

```
hist(data$lwage, 50, ylim = c(0, 350))
```



```
summary(data$jc)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.3389 0.00000 3.8330
```

```
print(quantile(data$jc, probs = c(0.01, 0.05, 0.1,
  0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%      5%     10%     25%     50%     75%     90%     95%
## 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 1.766667 2.266667
##      99%     100%
## 3.089665 3.833333
```

```
hist(data$jc, 50)
```

Histogram of data\$jc



```
summary(data$univ)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.000   0.000   0.200   1.926   4.200   7.500
```

```
print(quantile(data$univ, probs = c(0.01, 0.05, 0.1,
    0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##           1%           5%           10%           25%           50%           75%           90%
## 0.0000000 0.0000000 0.0000000 0.0000000 0.1999997 4.1999998 5.1777687
##           95%           99%          100%
## 5.9099934 7.5000000 7.5000000
```

```
hist(data$univ, 50, xlim = c(0, 8))
```

Histogram of data\$univ



```
summary(data$exper)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##       3.0   104.0   129.0   122.4   149.0   166.0
```

```
print(quantile(data$exper, probs = c(0.01, 0.05, 0.1,
    0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%    5%   10%   25%   50%   75%   90%   95%   99%  100%
##      25    56    74   104   129   149   160   163   166   166
```

```
hist(data$exper, 50, xlim = c(0, 200))
```

Histogram of data\$exper



```
summary(data$black)
```

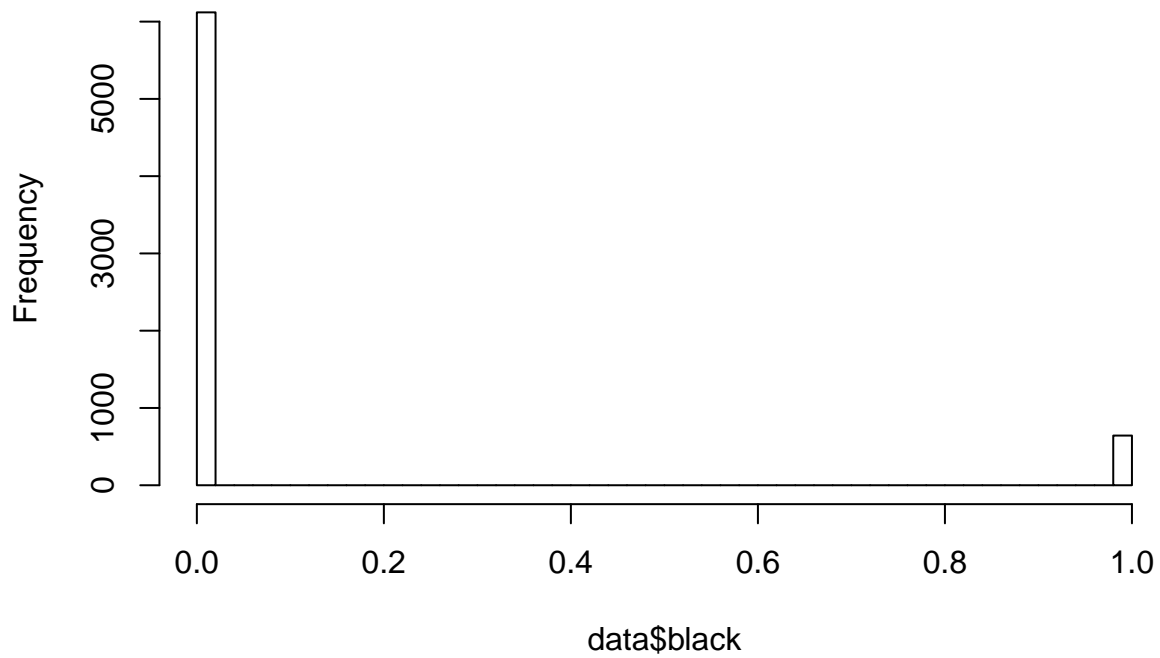
```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.09508 0.00000 1.00000
```

```
print(quantile(data$black, probs = c(0.01, 0.05, 0.1,
    0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%      5%     10%    25%    50%    75%    90%    95%    99%   100%
##       0       0       0       0       0       0       0       1       1       1
```

```
hist(data$black, 50)
```

Histogram of data\$black



```
summary(data$hispanic)
```

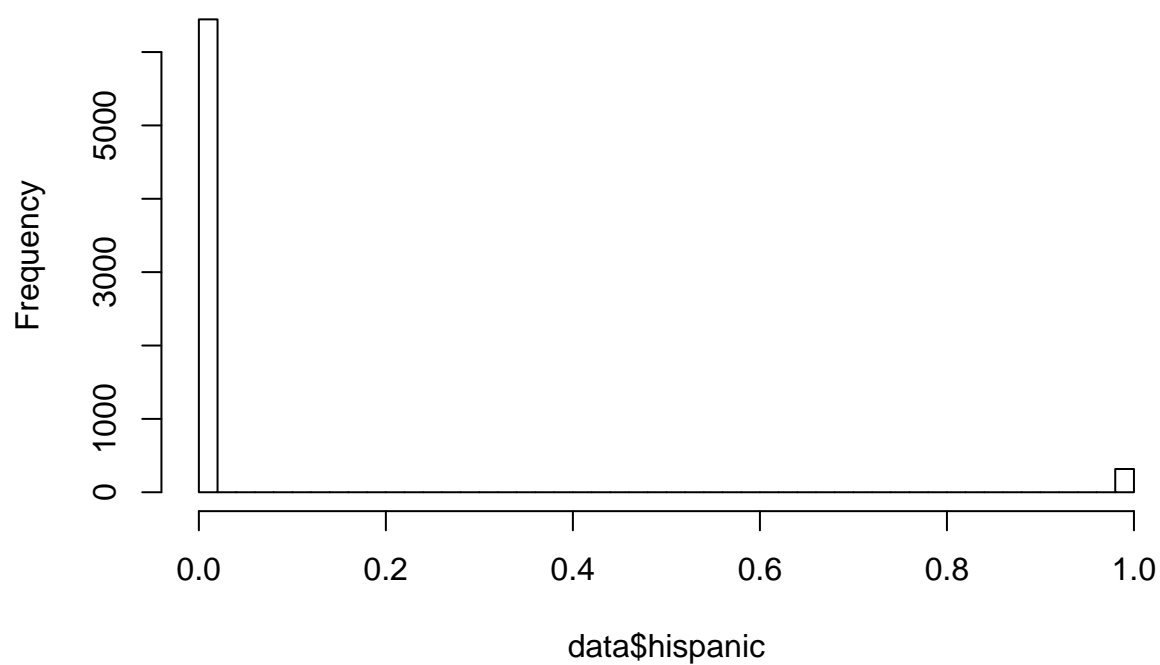
```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.04687 0.00000 1.00000
```

```
print(quantile(data$hispanic, probs = c(0.01, 0.05,
    0.1, 0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%      5%     10%     25%     50%     75%     90%     95%     99%    100%
##       0       0       0       0       0       0       0       0       1       1
```

```
hist(data$hispanic, 50)
```

Histogram of data\$hispanic



```
summary(data$AA)
```

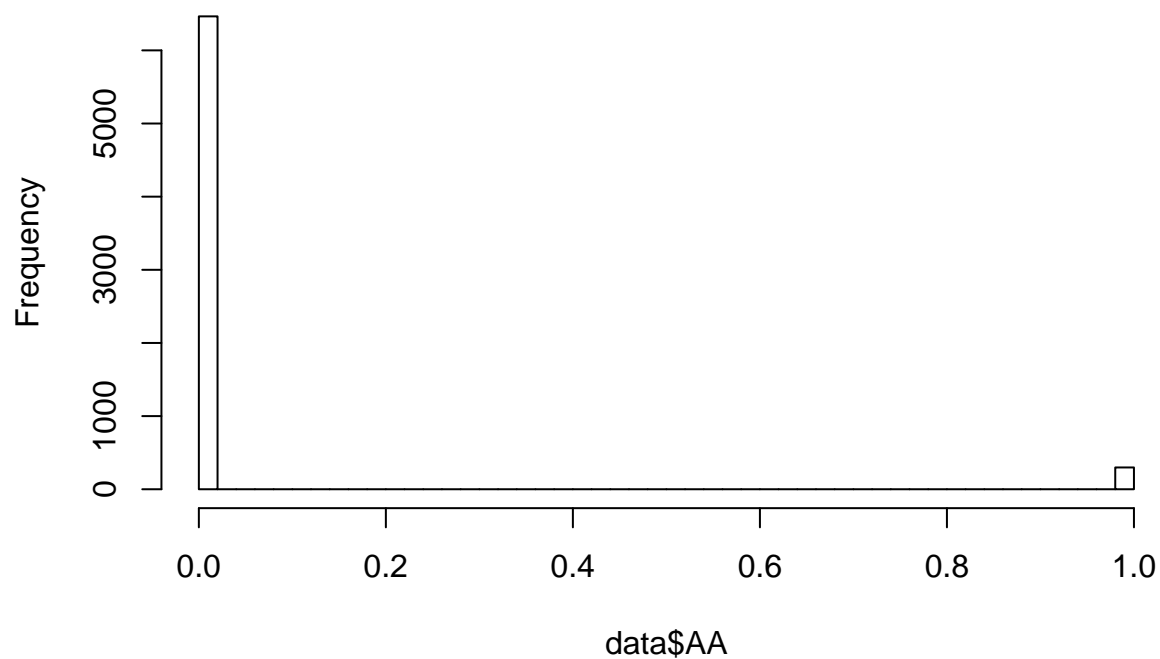
```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.00000 0.04406 0.00000 1.00000
```

```
print(quantile(data$AA, probs = c(0.01, 0.05, 0.1,
                                   0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%   5%  10%  25%  50%  75%  90%  95%  99% 100%
##       0    0   0   0   0   0   0   0   0   1   1
```

```
hist(data$AA, 50)
```

Histogram of data\$AA



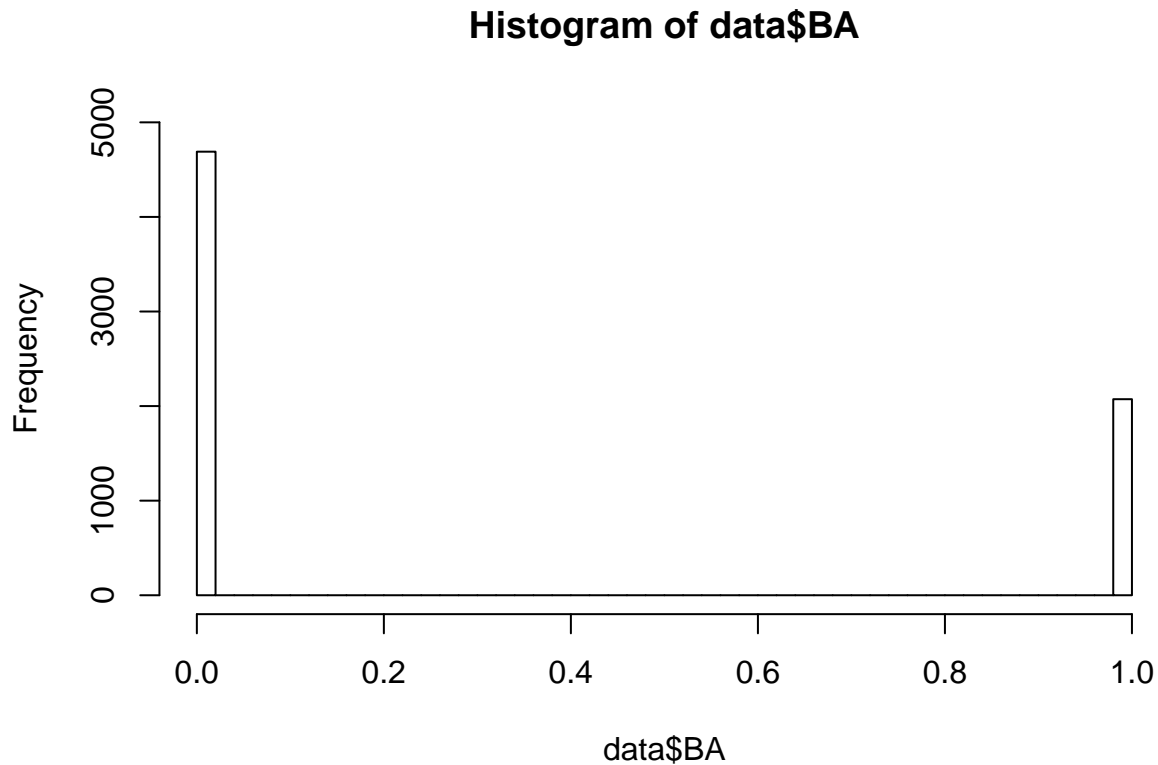
```
summary(data$BA)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.0000 0.0000 0.3065 1.0000 1.0000
```

```
print(quantile(data$BA, probs = c(0.01, 0.05, 0.1,
  0.25, 0.5, 0.75, 0.9, 0.95, 0.99, 1)))
```

```
##      1%   5%  10%  25%  50%  75%  90%  95%  99% 100%
##      0    0    0    0    0    1    1    1    1    1
```

```
hist(data$BA, 50, ylim = c(0, 5000))
```

Basic structure of the data

There are no missing values in the data.

lwage variable has a normal-like distribution.

jc variable has values from 0 to about 4 and is heavily positively skewed with a majority of values at or near 0.

univ variable has values from 0 to 7.5 and is heavily positively skewed with a majority of values at or near 0.
exper variable has values from 0 to 166 and is negatively skewed with a hill-climb distribution from 0 to about 500.

black, **hispanic**, **AA**, **BA** variables are binary with values of 0 or 1.


Question 2

```
# Create the experXblack variable by multiplying
# the exper and black variables.
data$experXblack = data$exper * data$black


# Run the requested OLS regression.
ols.lwage.8ind = lm(lwage ~ jc + univ + exper + black +
  hispanic + AA + BA + experXblack, data = data)
summary(ols.lwage.8ind)
```


```
##
## Call:
## lm(formula = lwage ~ jc + univ + exper + black + hispanic + AA +
##      BA + experXblack, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.11612 -0.27836  0.00432  0.28676  1.76811
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.4773315  0.0223780  66.017 < 2e-16 ***
## jc           0.0637926  0.0079034   8.072 8.15e-16 ***
## univ         0.0732806  0.0031486  23.274 < 2e-16 ***
## exper        0.0050234  0.0001667  30.141 < 2e-16 ***
## black        0.0331709  0.0613984   0.540  0.5890
## hispanic    -0.0193629  0.0248914  -0.778  0.4367
## AA          -0.0077759  0.0295497  -0.263  0.7924
## BA           0.0176735  0.0156553   1.129  0.2590
## experXblack -0.0012679  0.0004991  -2.541  0.0111 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4287 on 6754 degrees of freedom
## Multiple R-squared:  0.2282, Adjusted R-squared:  0.2272
## F-statistic: 249.6 on 8 and 6754 DF,  p-value: < 2.2e-16

# Print the diagnostic plots
plot(ols.lwage.8ind)
```




$\text{lm}(\text{lwage} \sim \text{jc} + \text{univ} + \text{exper} + \text{black} + \text{hispanic} + \text{AA} + \text{BA} + \text{experXblack})$





Fitted values
 $\text{lm}(\text{lwage} \sim \text{jc} + \text{univ} + \text{exper} + \text{black} + \text{hispanic} + \text{AA} + \text{BA} + \text{experXblack})$



```
# Print the B_hat4 and B_hat8 coefficients
print(ols.lwage.8ind$coefficients["black"])
```

```
##      black
## 0.03317088
```

```
print(ols.lwage.8ind$coefficients["experXblack"])
```

```
##  experXblack
## -0.001267898
```

Interpret the coefficients $\hat{\beta}_4$ and $\hat{\beta}_8$

$\hat{\beta}_4$ is the estimate for the black variable coefficient.

$\hat{\beta}_8$ is the estimate for the experXblack variable.

The β_4 coefficient captures the effect of being black on the log of wage, holding all other variables in the model fix and assuming zero-employment experience. It provides an indication of how much percentage point change to expect in the wage of an individual when they move from the baseline (white) to being black. Hence the coefficient captures the difference at the intercept of the log(wage) vs experience plot between black and white respondents.

The β_8 coefficient captures the effect of being black on the impact of experience over the years on wage. It is better explained in terms of derivatives with:

$$\frac{\delta \log(wage)}{\delta exper} = \beta_3 + \beta_8 * black$$

In the previous formulation, we can see that the β_8 coefficient captures the impact of being black on the slope of the log(wage) vs experience curve. In other words, the coefficient describes how much more or less experience impacts the log of wage for black people over the years, vs the baseline (white people). And because the outcome variable is the log of wage, the impact above can actually be formulated in terms of impact of ethnicity on percentage changes on the actual wage of individuals.

Question 3

```
# Show the summary of the model again
summary(ols.lwage.8ind)
```

```
##
## Call:
## lm(formula = lwage ~ jc + univ + exper + black + hispanic + AA +
##     BA + experXblack, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.11612 -0.27836  0.00432  0.28676  1.76811
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.4773315   0.0223780   66.017 < 2e-16 ***
## jc           0.0637926   0.0079034    8.072 8.15e-16 ***
## univ         0.0732806   0.0031486   23.274 < 2e-16 ***
## exper       0.0050234   0.0001667   30.141 < 2e-16 ***
## black       0.0331709   0.0613984    0.540  0.5890
## hispanic    -0.0193629   0.0248914   -0.778  0.4367
## AA         -0.0077759   0.0295497   -0.263  0.7924
## BA          0.0176735   0.0156553    1.129  0.2590
## experXblack -0.0012679   0.0004991   -2.541  0.0111 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4287 on 6754 degrees of freedom
## Multiple R-squared:  0.2282, Adjusted R-squared:  0.2272
## F-statistic: 249.6 on 8 and 6754 DF,  p-value: < 2.2e-16
```

```
# Print the univ coefficient
print(ols.lwage.8ind$coefficients["univ"])
```

```
##      univ
## 0.07328063
```

Test that the return to university education is 7%.

Null Hypothesis: $H_0: \beta_2 = 0.07$.

Alternate Hypothesis: $H_1: \beta_2 \neq 0.07$.

Using the linear model and summary statistics, we compute the pvalue for the test above as:

```
pvalue = 2 * (1 - pt((ols.lwage.8ind$coefficients["univ"] -  
    0.07)/coef(summary(ols.lwage.8ind))[, "Std. Error"] ["univ"],  
    df = summary(ols.lwage.8ind)$df[1]))  
pvalue
```

```
##      univ  
## 0.3246284
```

Based on the p-value, the test is not significant at the 0.05% significance level. Therefore, we can't reject the null hypothesis that the return to university education is 7%.

Question 4

Test that the return to junior college education is equal for black and non-black

From the linear model, we derive: $\frac{\delta \log(wage)}{\delta jc} = \beta_1$

The model is specified in a way that the return to junior college education, is β_1 , and is independent of ethnicity. Therefore without additional computation, we can immediately answer that the return on junior college education is the same for all ethnicities.

Question 5

Test whether the return to university education is equal to the return to 1 year of working experience.

Original model:

$$lwage = \beta_0 + \beta_1 jc + \beta_2 univ + \beta_3 exper + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \epsilon$$

Convert the experience variable from months to years by creating a new variable `experYr` that divides the original variable `exper` by 12. Replace the `exper` variable in the original model with this variable.

$$lwage = \beta_0 + \beta_1 jc + \beta_2 univ + \beta_3 experYr + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \epsilon$$

We would like to know if the β_2 and β_3 coefficients are the same or, equivalently, if their difference is 0. We can define a variable θ such that $\theta = \beta_2 - \beta_3$ and rewrite our model like this:

$$lwage = \beta_0 + \beta_1 jc + (\theta + \beta_3) univ + \beta_3 experYr + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \epsilon$$

Rewrite the model to get θ by itself as a coefficient:

$$lwage = \beta_0 + \beta_1 jc + \theta univ + \beta_3 (univ + experYr) + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \epsilon$$

Now our null hypothesis is $H_0: \theta = 0$.

Alternate Hypothesis: $H_1: \theta \neq 0$.


```

# Convert the exper variable from months to years
# by dividing it by 12.
data$experYr = data$exper/12
# Create a variable that is the sum of the univ and
# experYr variables
data$univ_plus_experYr = data$univ + data$experYr
# Rerun the regression with the new variables.
ols.lwage.univ.experYr = lm(lwage ~ jc + univ + univ_plus_experYr +
    black + hispanic + AA + BA + experXblack, data = data)
# Display a summary of the new model
summary(ols.lwage.univ.experYr)

```

```

##
## Call:
## lm(formula = lwage ~ jc + univ + univ_plus_experYr + black +
##     hispanic + AA + BA + experXblack, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.11612 -0.27836  0.00432  0.28676  1.76811
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.4773315   0.0223780   66.017 < 2e-16 ***
## jc              0.0637926   0.0079034    8.072 8.15e-16 ***
## univ            0.0129997   0.0035721    3.639 0.000276 ***
## univ_plus_experYr 0.0602810   0.0020000   30.141 < 2e-16 ***
## black           0.0331709   0.0613984    0.540 0.589038
## hispanic        -0.0193629   0.0248914   -0.778 0.436659
## AA              -0.0077759   0.0295497   -0.263 0.792446
## BA              0.0176735   0.0156553    1.129 0.258972
## experXblack      -0.0012679   0.0004991   -2.541 0.011088 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4287 on 6754 degrees of freedom
## Multiple R-squared:  0.2282, Adjusted R-squared:  0.2272
## F-statistic: 249.6 on 8 and 6754 DF,  p-value: < 2.2e-16

```

Based on the very low p-value (0.000276) for θ , the test is significant at the 0.05% significance level. And even though the value of θ is close to 0 at 0.0129997, we can reject the null hypothesis that $\theta = 0$.

Question 6

```

# Show the summary of the model again
summary(ols.lwage.8ind)

```

```

##
## Call:
## lm(formula = lwage ~ jc + univ + exper + black + hispanic + AA +

```

```
##      BA + experXblack, data = data)
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -2.11612 -0.27836  0.00432  0.28676  1.76811
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.4773315   0.0223780  66.017 < 2e-16 ***
## jc           0.0637926   0.0079034   8.072 8.15e-16 ***
## univ         0.0732806   0.0031486  23.274 < 2e-16 ***
## exper        0.0050234   0.0001667  30.141 < 2e-16 ***
## black        0.0331709   0.0613984   0.540  0.5890
## hispanic     -0.0193629   0.0248914  -0.778  0.4367
## AA           -0.0077759   0.0295497  -0.263  0.7924
## BA           0.0176735   0.0156553   1.129  0.2590
## experXblack -0.0012679   0.0004991  -2.541  0.0111 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4287 on 6754 degrees of freedom
## Multiple R-squared:  0.2282, Adjusted R-squared:  0.2272
## F-statistic: 249.6 on 8 and 6754 DF, p-value: < 2.2e-16
```

Test the overall significance of this regression.

We are testing the overall significance of the original model as stated below:

$$lwage = \beta_0 + \beta_1 jc + \beta_2 univ + \beta_3 exper + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \epsilon$$

Here is the output from the summary of the original model.

Residual standard error: 0.4287 on 6754 degrees of freedom

Multiple R-squared: 0.2282, Adjusted R-squared: 0.2272

F-statistic: 249.6 on 8 and 6754 DF, p-value: < 2.2e-16

1. Our model null hypothesis is that there is no relationship among any of the independent variables and lwage variable. We are able to reject the null hypothesis since our p-value of the f-statistic of the model is significant at < 2.2e-16.
2. Practical significance: we have an R-squared value of 0.2282, indicating that 22.82% of the variation in lwage is explained by our model.

Question 7

```
# Define a square term for the exper variable
data$experXexper = data$exper * data$exper
# Add the new variable to the regression
ols.lwage.9ind = lm(lwage ~ jc + univ + exper + black +
  hispanic + AA + BA + experXblack + experXexper,
```

```

data = data)
# Show the summary of the model
summary(ols.lwage.9ind)

##
## Call:
## lm(formula = lwage ~ jc + univ + exper + black + hispanic + AA +
##     BA + experXblack + experXexper, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.11982 -0.27743  0.00475  0.28741  1.77397
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.510e+00  4.427e-02  34.108 < 2e-16 ***
## jc           6.417e-02  7.916e-03   8.106 6.14e-16 ***
## univ         7.382e-02  3.211e-03  22.992 < 2e-16 ***
## exper        4.301e-03  8.588e-04   5.008 5.64e-07 ***
## black        2.994e-02  6.152e-02   0.487  0.6265
## hispanic     -1.932e-02  2.489e-02  -0.776  0.4378
## AA           -7.539e-03  2.955e-02  -0.255  0.7986
## BA           1.797e-02  1.566e-02   1.147  0.2513
## experXblack  -1.239e-03  5.002e-04  -2.477  0.0133 *
## experXexper   3.379e-06  3.939e-06   0.858  0.3911
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4287 on 6753 degrees of freedom
## Multiple R-squared:  0.2282, Adjusted R-squared:  0.2272
## F-statistic: 221.9 on 9 and 6753 DF,  p-value: < 2.2e-16

```

Estimated return to work experience in this model

$$lwage = \beta_0 + \beta_1 jc + \beta_2 univ + \beta_3 exper + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \beta_9 experXexper$$

We obtain the return on a year of experience by evaluating:

$$\frac{\delta lwage}{\delta exper} = \beta_3 + \beta_8 black + 2 * \beta_9 exper$$

$$\beta_3 = 4.301e - 03 = .004301, \beta_8 = .001239, \beta_9 = .000003379$$

Substituting these values in the the equation above we get:

$$\frac{\delta lwage}{\delta exper} = (.004301 - .001239 * black + .000006758 * exper)$$

Thus for blacks, the return to one year of working experience is:

```
.004301 - .001239 + .000006758
```

```
## [1] 0.003068758
```

We interpret is as 3/10th of a percent of increase in wage per year of experience.

And for whites, that return is:


```
.004301 + .000006758
```


```
## [1] 0.004307758
```

We interpret is as 4.3/10th of a percent of increase in wage per year of experience.


Question 8

```
# Plot the graphs from the model  
plot(ols.lwage.9ind)
```






Im(lwage ~ jc + univ + exper + black + hispanic + AA + BA + experXblack + e ...



Fitted values
 $\text{lm}(\text{lwage} \sim \text{jc} + \text{univ} + \text{exper} + \text{black} + \text{hispanic} + \text{AA} + \text{BA} + \text{experXblack} + \text{e} \dots)$



lm(lwage ~ jc + univ + exper + black + hispanic + AA + BA + experXblack + e ...

```
# Show the summary of the model
summary(ols.lwage.9ind)
```

```
##
## Call:
## lm(formula = lwage ~ jc + univ + exper + black + hispanic + AA +
##     BA + experXblack + experXexper, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.11982 -0.27743  0.00475  0.28741  1.77397
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.510e+00  4.427e-02  34.108  < 2e-16 ***
## jc           6.417e-02  7.916e-03   8.106  6.14e-16 ***
## univ         7.382e-02  3.211e-03  22.992  < 2e-16 ***
## exper        4.301e-03  8.588e-04   5.008  5.64e-07 ***
## black        2.994e-02  6.152e-02   0.487   0.6265
## hispanic     -1.932e-02  2.489e-02  -0.776   0.4378
## AA           -7.539e-03  2.955e-02  -0.255   0.7986
## BA           1.797e-02  1.566e-02   1.147   0.2513
## experXblack  -1.239e-03  5.002e-04  -2.477   0.0133 *
## experXexper   3.379e-06  3.939e-06   0.858   0.3911
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.4287 on 6753 degrees of freedom
## Multiple R-squared: 0.2282, Adjusted R-squared: 0.2272
## F-statistic: 221.9 on 9 and 6753 DF, p-value: < 2.2e-16

# Use the robust standard errors
coeftest(ols.lwage.9ind, vcov = vcovHC)

##
## t test of coefficients:
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.5101e+00 4.3591e-02 34.6427 < 2.2e-16 ***
## jc          6.4168e-02 7.6224e-03  8.4183 < 2.2e-16 ***
## univ        7.3819e-02 3.4501e-03 21.3963 < 2.2e-16 ***
## exper       4.3008e-03 8.4541e-04  5.0873 3.731e-07 ***
## black       2.9937e-02 6.8436e-02  0.4374 0.66180
## hispanic    -1.9317e-02 2.4985e-02 -0.7731 0.43947
## AA          -7.5392e-03 2.7481e-02 -0.2743 0.78383
## BA          1.7967e-02 1.6579e-02  1.0837 0.27853
## experXblack -1.2388e-03 5.3539e-04 -2.3139 0.02071 *
## experXexper  3.3790e-06 3.8745e-06  0.8721 0.38318
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Homoskedasticity analysis:

We are testing homoskedasticity of the model with the quadratic experience term, expressed as:

$$lwage = \beta_0 + \beta_1 jc + \beta_2 univ + \beta_3 exper + \beta_4 black + \beta_5 hispanic + \beta_6 AA + \beta_7 BA + \beta_8 experXblack + \beta_9 experXexper + \epsilon$$

We observe a small amount of heteroskedasticity from the plot:

- 1 - We can see from the residuals vs fitted plot that the variance band changes very slightly as we move to higher fitted values.
- 2 - The same story is told by the scale-location plot where we see that the smoothing line is not quite completely horizontal. Therefore we conclude a very small amount of heteroskedasticity and decide to use heteroskedasticity-robust methods for coefficient estimation.
- 3 - We do not look at the Breusch Pagan test since we have a large number of observations, therefore we know almost certainly that we will obtain significance.

The implication of heteroskedasticity (even small) in the data is that the standard error of the univ coefficient (β_2) may be biased. Biased standard errors can impact the outcomes of statistical tests. Therefore, it can affect the testing of no effect of university education on salary, which is the t-test on the coefficient β_2 .

The β_2 coefficient from the robust method was essentially unchanged, going from a value of $\beta_2 = 7.382e - 02$ using the non-robust estimation to $\beta_2 = 7.3819e - 02$ with robust estimation. The standard error of the β_2 coefficient was changed, going from $3.211e-03$ using the non-robust estimation to $3.4501e-03$ using the robust estimation. However, even using the robust estimation, the p-value for our β_2 coefficient remains significant at the 0.05, and we confirm that we must reject the null hypothesis that the coefficient value is null, with the meaning of that null hypothesis being that there is no relationship between time at university and wages.